# Online Result Summary

## Model: text_reg_batch

GPU(s):

Total Available GPU Memory: 0 GB

Constraint targets: None

In 53 measurements across 9 configurations, **text_reg_batch_config_4** is **50%** better than the default configuration at maximizing throughput, under the given constraints, on GPU(s) .

- **text_reg_batch_config_4**: 0 GPU instance with a max batch size of 16 on platform onnxruntime

Curves corresponding to the 3 best model configuration(s) out of a total of 9 are shown in the plots.

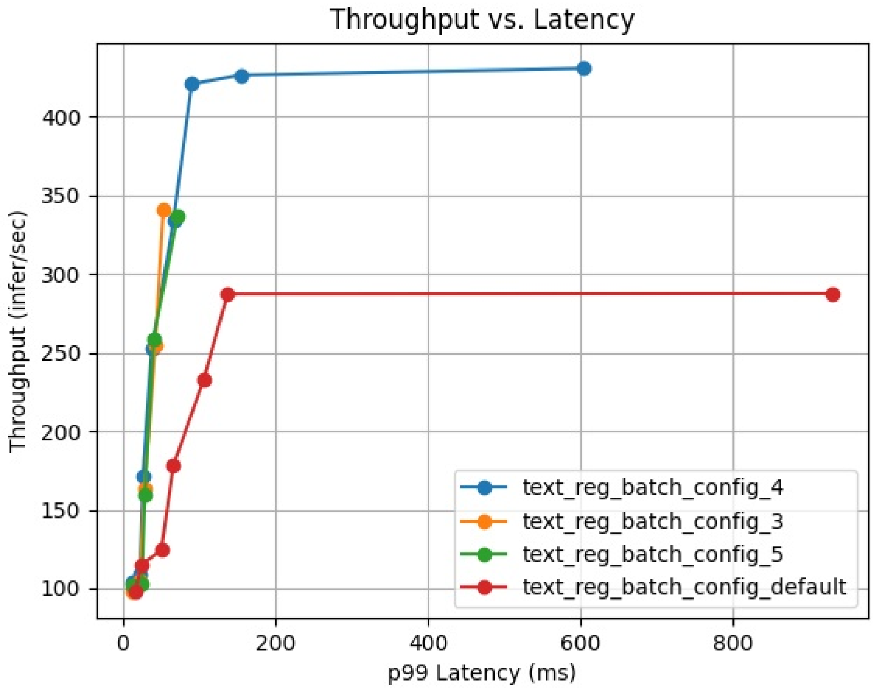

**Throughput vs. Latency curves for 3 best configurations.**

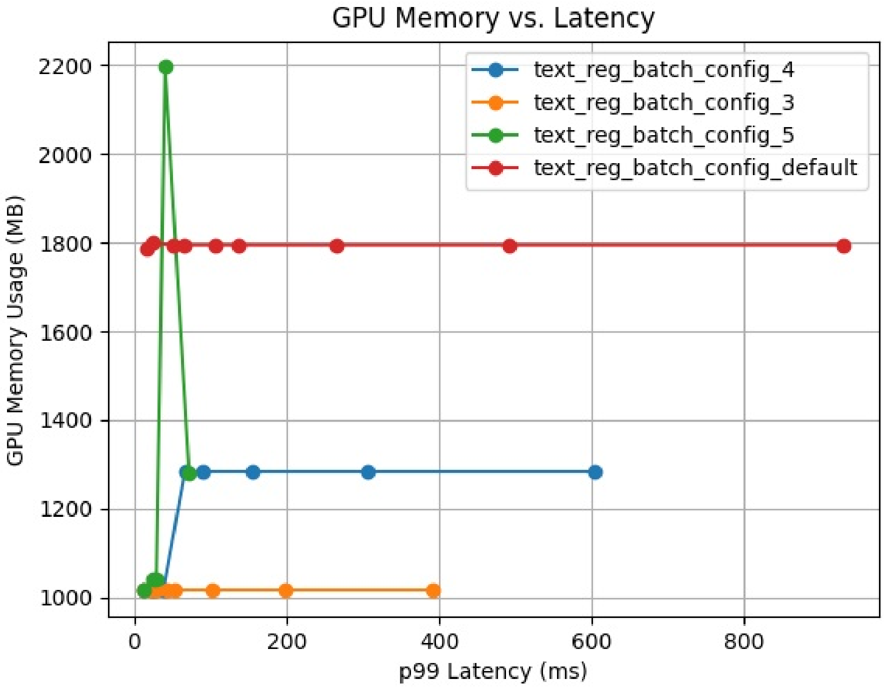

**GPU Memory vs. Latency curves for 3 best configurations.**

The following table summarizes each configuration at the measurement that optimizes the desired metrics under the given constraints.

| **Model Config Name** | **Max Batch Size** | **Dynamic Batching** | **Total Instance Count** | **p99 Latency (ms)** | **Throughput (infer/sec)** | **Max GPU Memory Usage (MB)** | **Average GPU Utilization (%)** |
|---|---|---|---|---|---|---|---|
| text_reg_batch_config_4 | 16 | Enabled | 0:GPU | 603.053 | 430.757 | 1283 | 94.5 |
| text_reg_batch_config_3 | 8 | Enabled | 0:GPU | 53.717 | 341.062 | 1015 | 93.5 |
| text_reg_batch_config_5 | 32 | Enabled | 0:GPU | 72.569 | 336.716 | 1279 | 87.7 |
| text_reg_batch_config_default | 8 | Enabled | 2:GPU | 930.668 | 287.462 | 1795 | 74.0 |